

Learn Latent Representations

VLC = $\int \int \int$ Vision
Learning & Control

Unsupervised and self-supervised learning
Auto-Encoders

Jonathon Hare


Vision, Learning and Control
University of Southampton

Compressed and Latent Representations


Compression

Autoencoders

The aim of an autoencoder is to learn a representation (encoding) for a set of data, typically for dimensionality reduction, by training the network to ignore signal noise.



Deep Autoencoders



1

¹Image taken from wikipedia

Jonathon Hare

Auto-encoders

7 / 11

Denoising Autoencoders

- Denoising autoencoders take a partially corrupted input and train to recover the original undistorted input.
- To train an autoencoder to denoise data, it is necessary to perform a preliminary stochastic mapping to corrupt the data ($x \rightarrow \tilde{x}$).
- A normal autoencoder is used with \tilde{x} is used as input and x as output.
- In a denoising autoencoder, the loss should be computed on $\mathcal{L}(x, \hat{x})$ as opposed to $\mathcal{L}(\tilde{x}, \hat{x})$.

Sparse Autoencoders

- In a sparse autoencoder, there are more hidden units than inputs, but only a small number of the hidden units are allowed to be active at the same time.

Convolutional Autoencoders

-

- Any basic AE (or its variant) is used to learn a compact representation of data.
- You can learn automatic features from data.
- Denoising can help generalise over the test set since the data is distorted by adding noise.
- Pretraining networks by learning your network weights using a stacked AE.